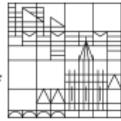


Chapter 1

Images, Noise and Filters

University of
Konstanz



**Lecture “Image Analysis and Computer Vision”
Winter semester 2014/15
Bastian Goldlücke**

1 Image types and discretization

2 Image noise

3 Correlation and Convolution

Noise reduction and moving average

The correlation

Gaussian filter

The convolution

Other filter kernels

4 Non-linear filters and denoising

Bilateral filter

Non-local means

Median filter

5 Summary

1 Image types and discretization

2 Image noise

3 Correlation and Convolution

Noise reduction and moving average

The correlation

Gaussian filter

The convolution

Other filter kernels

4 Non-linear filters and denoising

Bilateral filter

Non-local means

Median filter

5 Summary

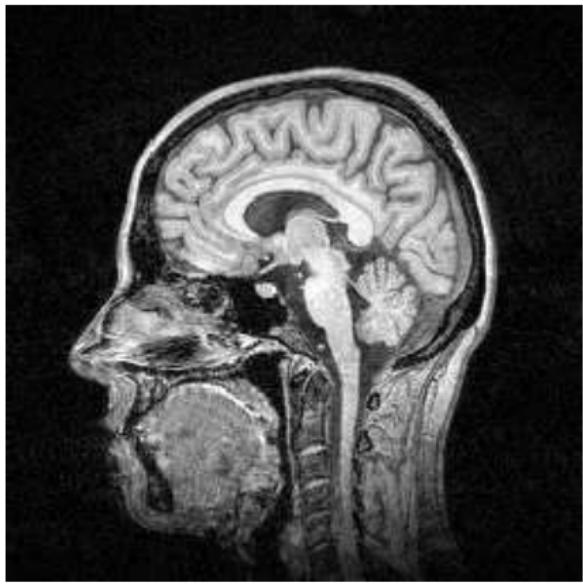
Continuous greyscale image

- mapping f from a rectangular domain $\Omega = (0, w) \times (0, h)$ to a co-domain \mathbb{R} ,

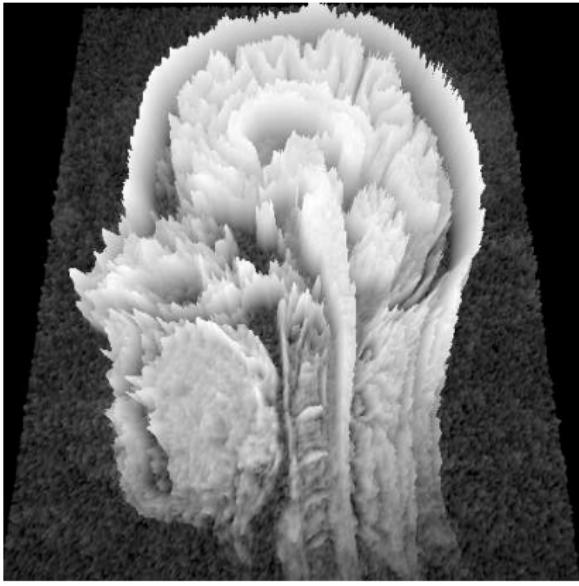
$$f : \mathbb{R}^2 \supset \Omega \rightarrow \mathbb{R}$$

- domain Ω is called *image domain* or *image plane*
- co-domain specifies grey value
- Usually low grey values mean dark and high grey values bright.

Continuous greyscale image: visualization



Magnetic resonance (MR) image of a human head



Representation as a function f over a rectangular image domain Ω

- discretisation of the **domain** Ω
- Image data are only given on a rectangular point grid within the image domain Ω
- creates a *digital image*

$$(f_{i,j})_{i=1,\dots,N, j=1,\dots,M}$$

- The grid point or cell (i, j) is called *pixel* (picture element).
- 2D images often have equal pixel distances in both directions.
- Image processing people often normalise these grid sizes to 1. This can be dangerous and is not recommended!
- If the sampling is too coarse, the image quality degrades severely.

Visualization of different sampling rates



256×256 pixels



128×128 pixels

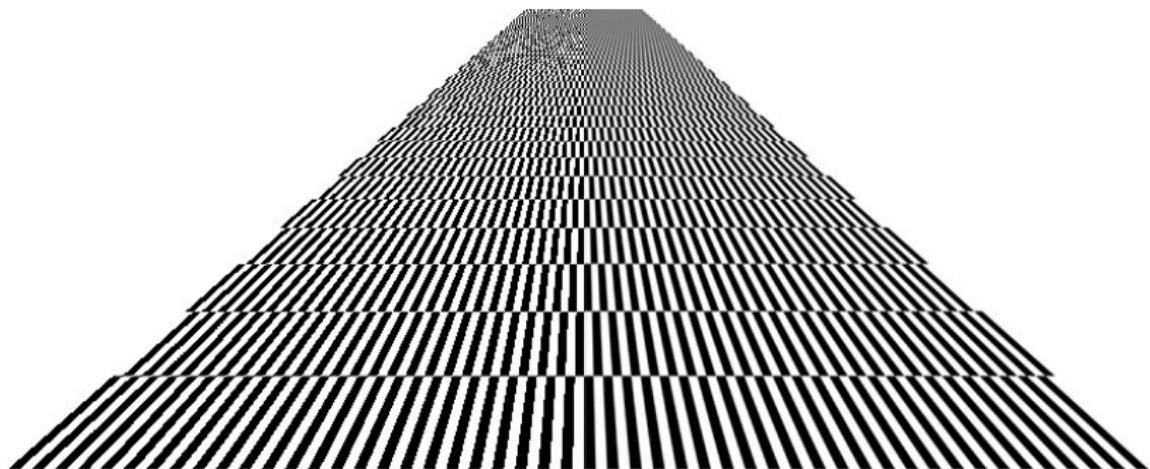


64×64 pixels



32×32 pixels

A sampling rate which is too low creates *aliasing* ...



- Aliasing means signals become indistinguishable when sampled
- Reminder (maybe) from signal processing: sampling rate must be above the *Nyquist rate*, which is two times the bandwidth of the signal, to have an aliasing-free result
- We will get back to that at a later point, when we discuss the Fourier transform of an image

- discretisation of the **co-domain**
- saves disk space
- example: if a grey value is coded by a single byte, the discrete co-domain is given by

$$\{0, 1, \dots, 255\}$$

- *Binary images* have the co-domain $\{0, 1\}$.
- Humans can distinguish only 40 greyscales.
- They are even very good in recognising the content of binary images.

Visualization of quantisation effects



256 greyscales



32 greyscales



8 greyscales



2 greyscales

- generalization: domain in \mathbb{R}^m
- $m = 1$: signals
- $m = 2$: (two-dimensional) images
- $m = 3$: three-dimensional images
 - important in medical imaging, e.g. computerised tomography (CT, Computertomographie), magnetic resonance imaging (MRI, Kernspintomographie).
 - image points/cells in 3-D are called *voxels* (volume elements).
 - voxel dimensions usually differ in different directions!

Rendering of an ultrasound image of a fetus in its 10th week

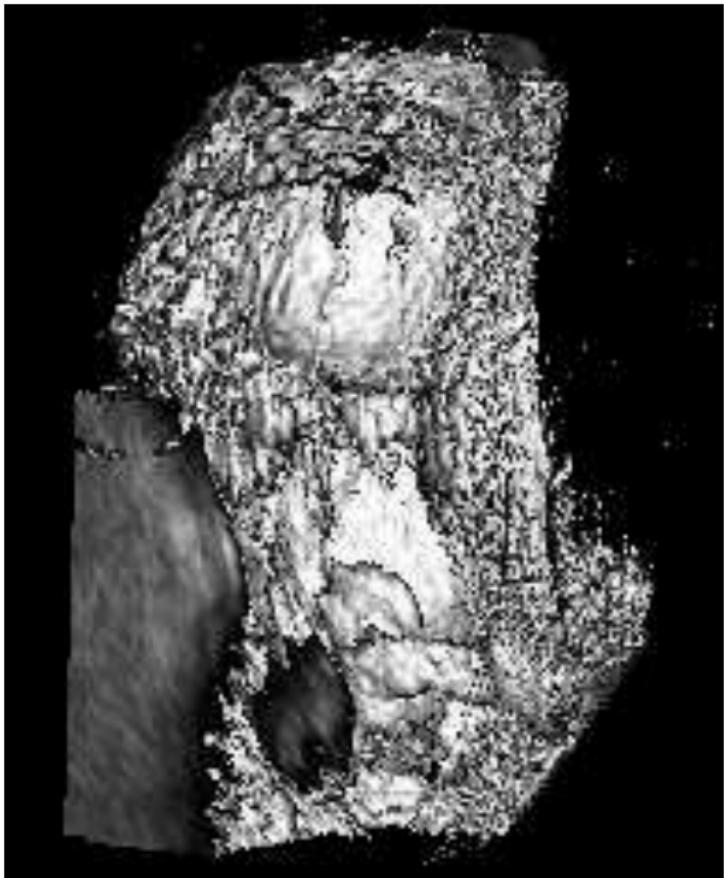
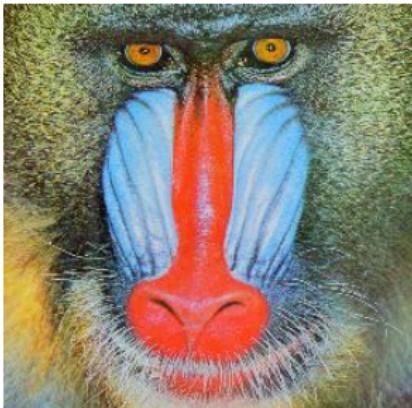


Image: J. Weickert, K. Zuiderveld, B.M. ter Haar Romeny, W. Niessen (1997).

Vector-valued images

- generalization: co-domain in \mathbb{R}^n , containing n channels
- Example 1: Colour images
 - Three channels: R (red), G (green), B (blue).
 - Humans can distinguish 2.000.000 colours!
 - We will learn much more about color perception in a later lecture
- Example 2: Multispectral image (e.g. satellite images)
 - numerous channels that represent different frequency bands

Color image as example for a vector-valued image



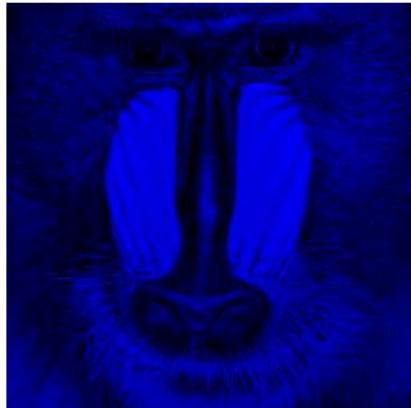
Original image



Red channel



Green channel

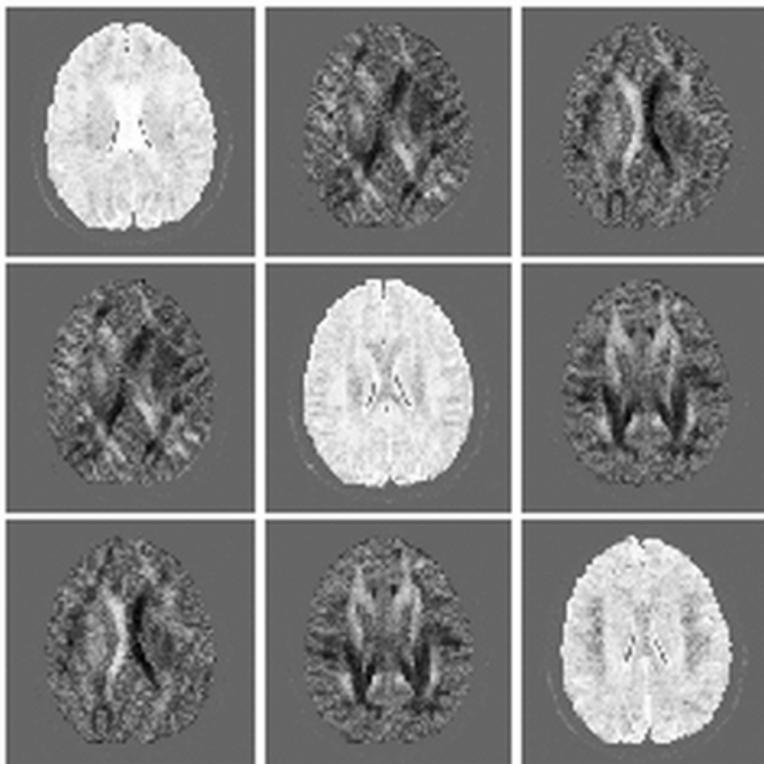


Blue channel

Matrix-valued images (sometimes: “tensor”)

- co-domain in $\mathbb{R}^{n \times n}$
- Example: diffusion tensor MRI (DT-MRI, “Diffusionstensor-Kernspintomographie”): measures in each voxel of a 3D image domain a symmetric positive definite 3×3 matrix
- may create additional constraints, e.g.: A reasonable filter should not destroy relevant properties such as positive definiteness of the matrix field.

Nine coefficients of a DT-MRI data set of a human brain



Since the diffusion matrix is a symmetric 3×3 matrix, only 6 out of 9 images differ.

Image sequences

- One can consider image sequences for any of the above mentioned image types.
- This increases the dimensionality of the domain from m to $m + 1$.
- Example 1: A color video sequence can be regarded as a 3D image with 3D co-domain.
- Example 2: 3D ultrasound can create a sequence of 3D scalar-valued images, which can be regarded as 4D image.

Four subsequent frames of an image sequence



Relevant images in this class

- We mainly focus on 2D scalar or RGB images ($m = 2$, $n = 1$ or $n = 3$).
- The gradient of a greyscale image yields another important vector-valued image ($m = 2$, $n = 2$).
- Tensors are lurking in the shadows of the second chapter.
- Many methods can also be generalised to images in higher dimensions (see the Mathematical Image Analysis lecture).

Notation and conventions: digital images

- A digital grayscale image will be usually denoted with

$$f = (f_{i,j}),$$

where we assume $f_{i,j} \in \mathbb{R}$.

- Pixel coordinates (i, j) can be anywhere in \mathbb{Z}^2 , but to make everything well-defined, we'll implicitly assume that an image always has “compact support”, i.e. $f_{i,j} = 0$ outside a bounded region $\Omega \subset \mathbb{Z}^2$ (the usually rectangular image domain).
- Sometimes I write $f(i, j)$ instead of $f_{i,j}$ if the formulas become more easy to read that way.
- If the image has multiple channels (i.e. RGB), I make that clear by e.g. writing

$$\mathbf{f} : \mathbb{Z}^2 \rightarrow \mathbb{R}^3$$

to indicate the co-domain. I also use a bold-faced letter \mathbf{f} .

Overview

1 Image types and discretization

2 Image noise

3 Correlation and Convolution

Noise reduction and moving average

The correlation

Gaussian filter

The convolution

Other filter kernels

4 Non-linear filters and denoising

Bilateral filter

Non-local means

Median filter

5 Summary

Image noise



Original image



Noisy image

- How can one characterize noise, what **types of noise** are there?
- How can one get rid of noise?

Additive noise

- Most important type of noise
- Grey values and noise are assumed to be independent:

$$f_{i,j} = g_{i,j} + n_{i,j}$$

where

$g = (g_{i,j})$ is the original image,
 $n = (n_{i,j})$ the noise, and
 $f = (f_{i,j})$ the noisy image.

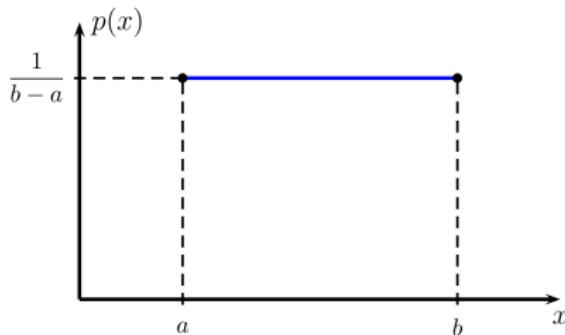
- The noise n may have different distributions, e.g.
 - uniform distribution (very simple)
 - Gaussian distribution (very common)

Uniform distribution

- often not the most realistic noise model, but easy to simulate
- has a constant density function within some interval $[a, b]$:

$$p(x) = \begin{cases} \frac{1}{b-a} & \text{for } x \in [a, b], \\ 0 & \text{else.} \end{cases}$$

- appears e.g. in connection with quantisation



Density function for uniform noise. Author: M. Mainberger (2008).

Uniform additive noise



Original image

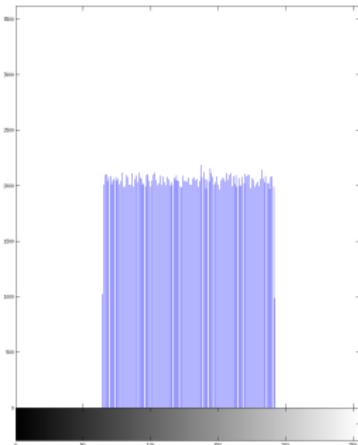
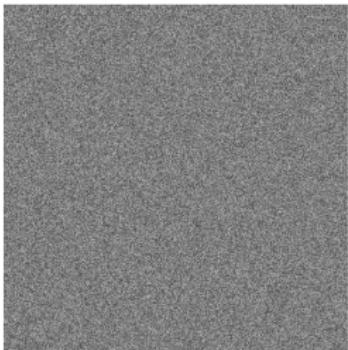


Noisy, $(b - a) = 0.5f_{max}$

Uniform additive noise

```
g = ones([512,512])*128;
a = -64; b = 64;
n = a + (b-a)*rand(size(image));
f = g + n;

figure,
subplot(1,2,1)
imshow(f / 255.0),
subplot(1,2,2)
imhist(uint8(f))
```



Noise function and histogram

Gaussian distribution

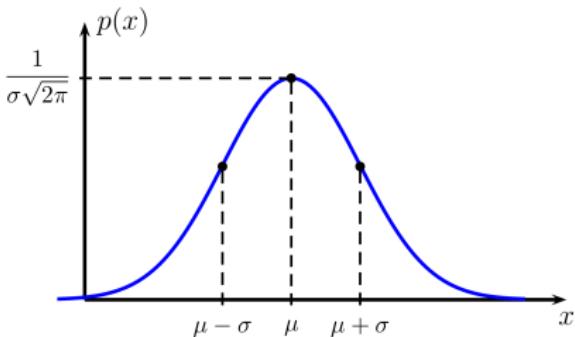
- most important noise model
- good approximation in many practical situations, e.g.
 - thermal sensor noise in CCD cameras
 - circuit noise caused by signal amplifications
- has density function

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right),$$

where μ is the mean and $\sigma > 0$ the standard deviation.

- Note: σ^2 is also called the “variance” of the distribution.

Gaussian distribution



Density function for Gaussian noise. Author: M. Mainberger (2008).

- For a Gaussian-distributed random variable X the following probabilities hold:

$$P(\mu - \sigma \leq X \leq \mu + \sigma) \approx 68\%$$

$$P(\mu - 2\sigma \leq X \leq \mu + 2\sigma) \approx 95.5\%$$

$$P(\mu - 3\sigma \leq X \leq \mu + 3\sigma) \approx 99.7\%$$

- Hence, Gaussian noise lives almost completely in the interval $[\mu - 3\sigma, \mu + 3\sigma]$.
- This is useful to know for discretization, since the distribution technically has infinite support.

Gaussian additive noise



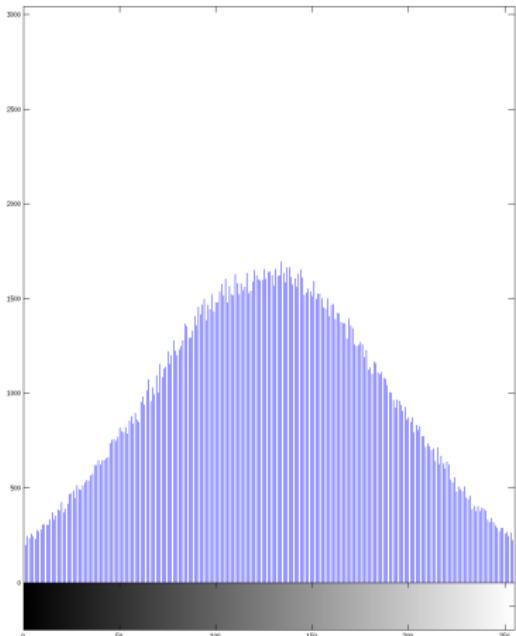
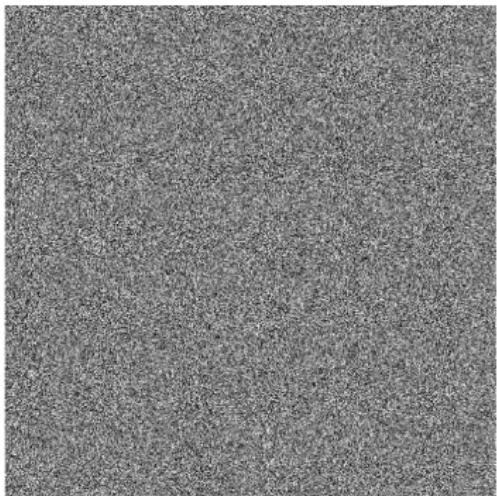
Original image



Additive Gaussian noise, $\sigma = 0.25f_{max}$

Gaussian additive noise

```
g = ones([512,512])*128;  
sigma=64;  
n = sigma*randn(size(image));  
f = g + n;
```



Noise function and histogram

Multiplicative noise

- Multiplicative noise is signal-dependent (unlike additive noise).
- Frequently, it is proportional to the grey value:

$$f_{i,j} = g_{i,j} + n_{i,j}g_{i,j} = (1 + n_{i,j})g_{i,j}.$$

- The noise n can again follow different distributions.
- Multiplicative noise is prevalent in radar, ultrasound and tomographic images.

Multiplicative Gaussian noise



Original image



Multiplicative Gaussian noise, $\sigma = 0.25f_{max}$

Impulse Noise

- degrades the image at **some (!)** pixels where erroneous grey values are created, in contrast to additive or multiplicative noise that affects all pixels
- Example: pixel defects in a CCD chip of a digital camera
- **Unipolar impulse noise** gives degradations having only one grey value, whereas **bipolar noise** attains two grey values.
- Bipolar noise with highest and lowest grey values is called **salt-and-pepper noise**.

Salt-and-pepper noise



Original image



Salt-and-pepper noise, 20 percent of pixels affected

We now have many ways to add interesting noise to images ...
but how to get rid of it again?

Overview

1 Image types and discretization

2 Image noise

3 Correlation and Convolution

Noise reduction and moving average

The correlation

Gaussian filter

The convolution

Other filter kernels

4 Non-linear filters and denoising

Bilateral filter

Non-local means

Median filter

5 Summary

Question: noise reduction

- Given a camera and a still scene, how can you reduce noise?

Question: noise reduction

- Given a camera and a still scene, how can you reduce noise?
- Take lots of images and average them!



Question: noise reduction

- Given a camera and a still scene, how can you reduce noise?
- Take lots of images and average them!



- What's the next best thing?

Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Some function



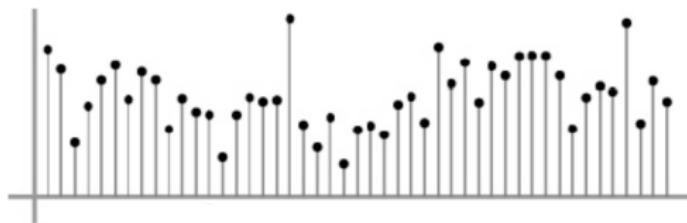
Local image data

		7

Modified image data

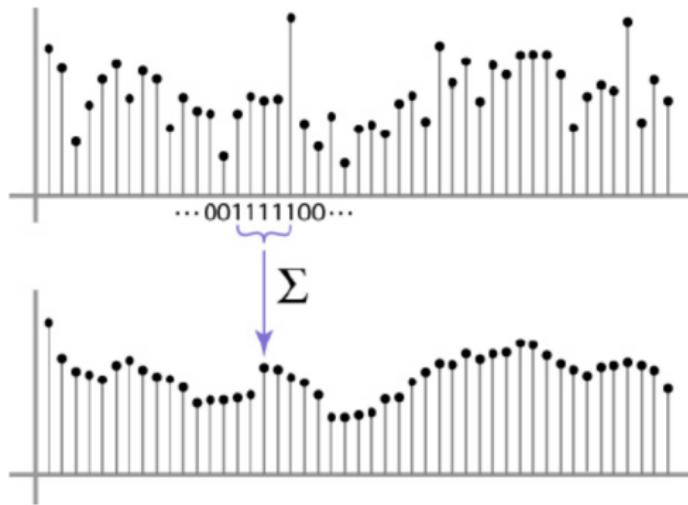
Moving average in 1D

- Simplest thing: replace each pixel by the average of its neighbors
- This assumes that neighboring pixels are similar, and the noise to be independent from pixel to pixel.



Moving average in 1D

- Simplest thing: replace each pixel by the average of its neighbors
- This assumes that neighboring pixels are similar, and the noise to be independent from pixel to pixel.
- “Kernel” for moving average in 1D: $\frac{[1,1,1,1,1]}{5}$.



Example: stock market indices



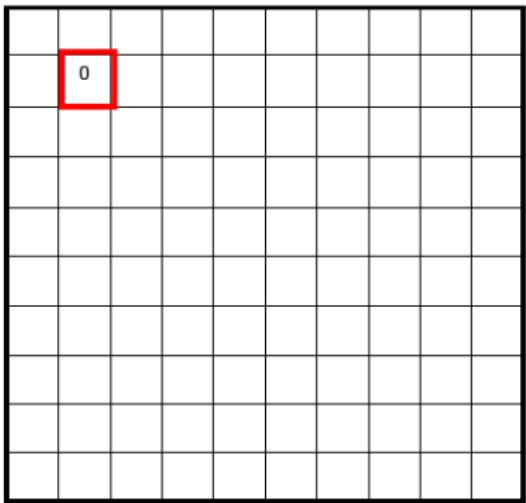
German stock market index (DAX) on April 17, 2014. **Blue:** Daily values. **Green:** Averaged over the last 38 days. **Red:** Averaged over the last 200 days.

Moving average in 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$



Moving average in 2D

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

0	10										

Moving average in 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10	20								

Moving average in 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving average in 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10	20	30	30	

Moving average in 2D

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Generalization: the correlation

- Involves weighted combinations of pixels in small neighborhoods.
- The output pixel's value is determined as a weighted sum of input pixel values

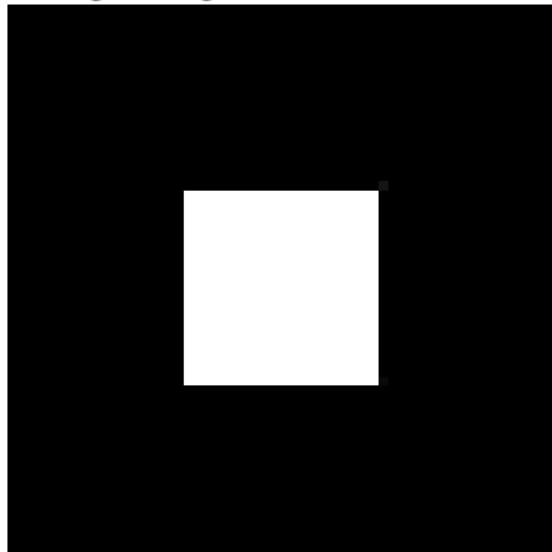
$$g(i, j) = \sum_{k, l \in \mathbb{Z}} f(i + k, j + l) h(k, l)$$

- The entries of the **weight kernel** or **mask** $h(k, l)$ are often called the filter coefficients.
- The operator defined in this way is the **correlation operator** and denoted as

$$g = f \otimes h.$$

Visualizing a 2D kernel

The 3×3 box filter:
moving average


$$\dots \quad k = 0 \quad \vdots \quad \dots \quad l = 0$$

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0

$$\vdots$$

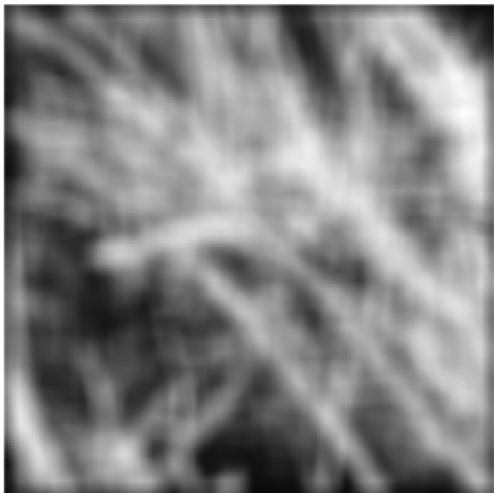
Smoothing with a box (local mean) filter



depicts box filter:
white = high value, black = low value



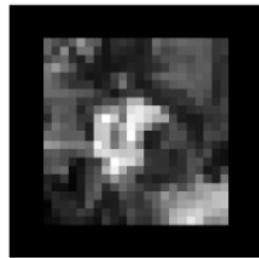
original



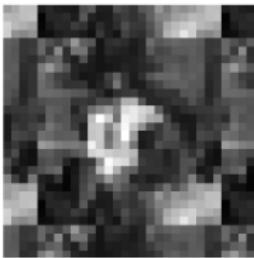
filtered

Boundary conditions

- The results of filtering the image in this form will lead to a darkening of the corner pixels.
- The original image is effectively being padded with 0 values wherever the convolution kernel extends beyond the original image boundaries.
- A number of alternative padding or extension modes have been developed.



zero



wrap



clamp



mirror

The Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?
- Removes high-frequency components from the image (low-pass filter).

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$F[x, y]$

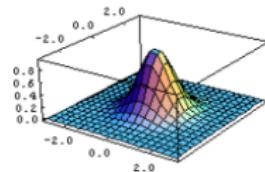
$$\frac{1}{16}$$

1	2	1
2	4	2
1	2	1

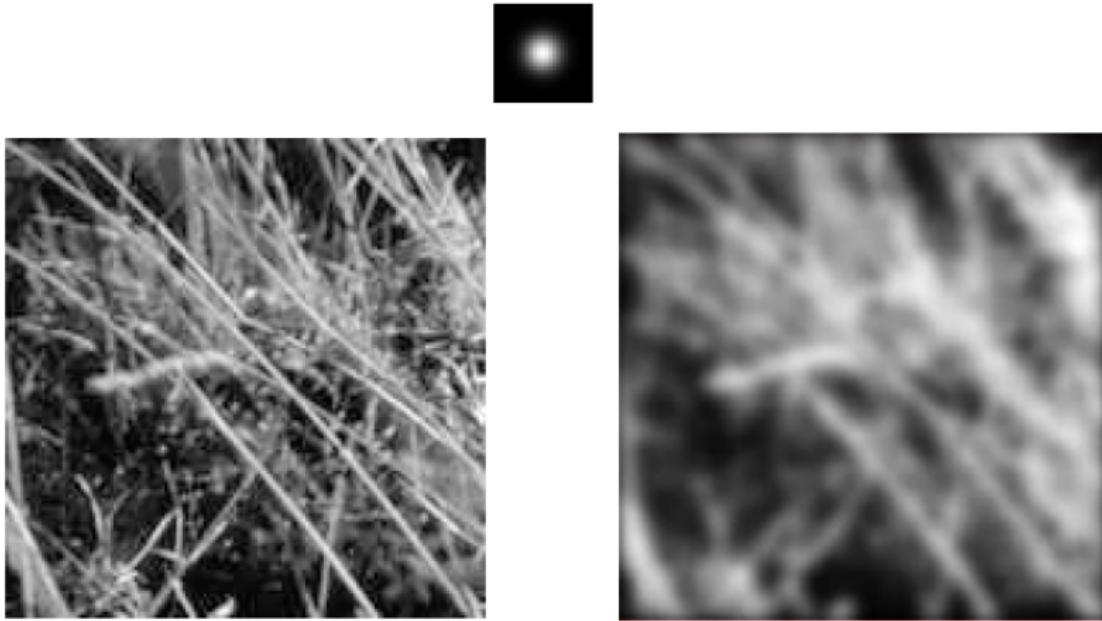
$$H[u, v]$$

This kernel is an approximation of a 2d Gaussian function:

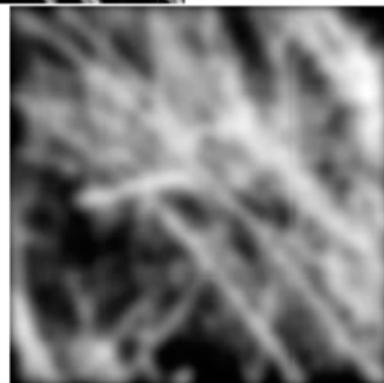
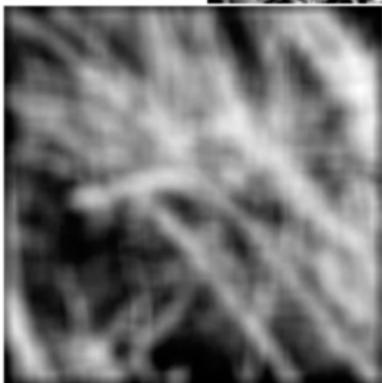
$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



Smoothing with a Gaussian filter

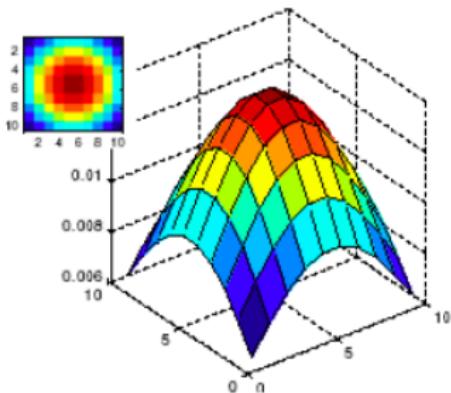


Mean vs. Gaussian

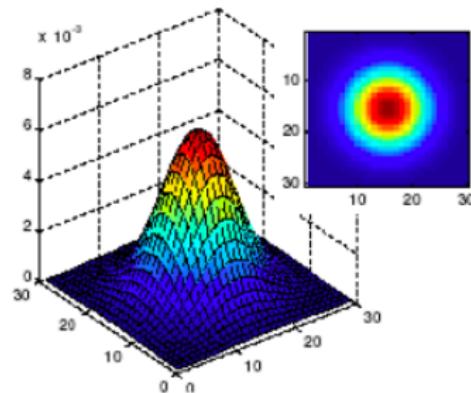


Gaussian filter parameters

- **Size of kernel or mask:** Gaussian function has infinite support, but discrete filters use finite kernels.



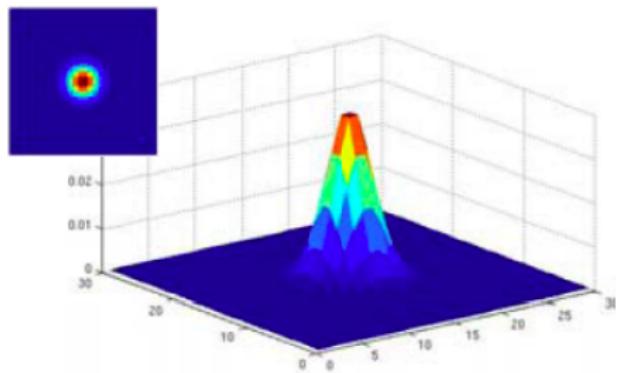
$\sigma = 5$ with
10 x 10
kernel



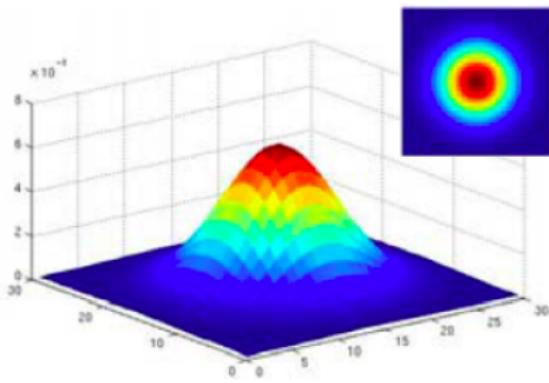
$\sigma = 5$ with
30 x 30
kernel

Gaussian filter parameters

- **Standard deviation of the Gaussian:** determines extent of smoothing.

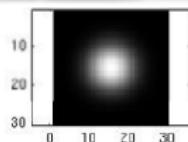
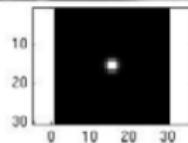


$\sigma = 2$ with
30 x 30
kernel

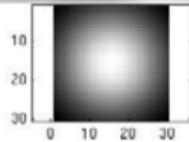


$\sigma = 5$ with
30 x 30
kernel

Gaussian filter parameters



...



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

Impulse response of correlation

- Impulse signal δ :

$$\delta_{i,j} = \begin{cases} 1 & \text{if } i = 0 \text{ and } j = 0, \\ 0 & \text{else.} \end{cases}$$

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G[x, y]$$

Impulse response of correlation

- Impulse signal δ :

$$\delta_{i,j} = \begin{cases} 1 & \text{if } i = 0 \text{ and } j = 0, \\ 0 & \text{else.} \end{cases}$$

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G[x, y]$$

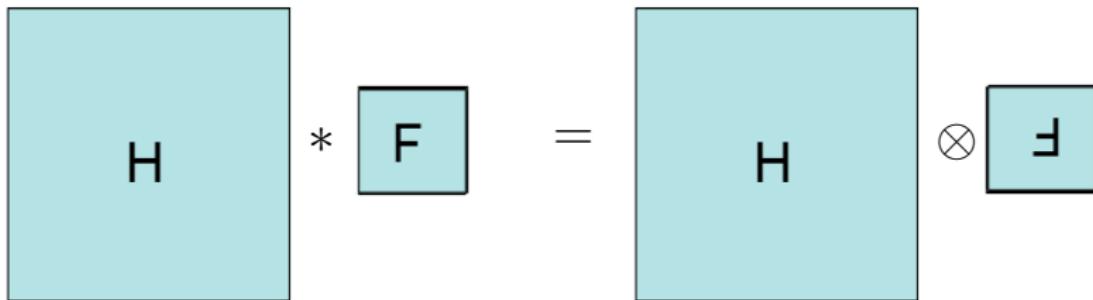
- The filter will be flipped, mathematicians do not like that ...
 - no left-side “identity element” for the correlation operation?
 - not commutative?

Convolution

- The convolution $g = f * h$ of two images f, h is defined as

$$g(i, j) = \sum_{k, l \in \mathbb{Z}} f(i - k, j - l) h(k, l).$$

- Since now $\delta * h = h$, the kernel h is sometimes called the **impulse response function**.
- Equivalent to flipping the filter in both dimensions (bottom to top, right to left) and apply cross-correlation.



Convolution vs. Correlation

- Convolution (Matlab: `conv2`)

$$g(i,j) = \sum_{k,l \in \mathbb{Z}} f(i-k, j-l) h(k, l)$$
$$g = f * h$$

- Correlation (Matlab: `imfilter`)

$$g(i,j) = \sum_{k,l \in \mathbb{Z}} f(i+k, j+l) h(k, l)$$
$$g = f \otimes h$$

- for filters which are symmetric in both dimensions, the output is the same (e.g. Gaussian, box filters)
- for correlation, visualization of the filter is much more intuitive

Convolution vs. Correlation

- Convolution (Matlab: `conv2`)

$$g(i,j) = \sum_{k,l \in \mathbb{Z}} f(i-k, j-l) h(k, l)$$
$$g = f * h$$

- Correlation (Matlab: `imfilter`)

$$g(i,j) = \sum_{k,l \in \mathbb{Z}} f(i+k, j+l) h(k, l)$$
$$g = f \otimes h$$

- for filters which are symmetric in both dimensions, the output is the same (e.g. Gaussian, box filters)
- for correlation, visualization of the filter is much more intuitive
- ... **that's why they are very frequently mixed up - be aware**

Useful properties of the convolution

- **Commutative:** $f * h = h * f$
- **Associative:** $(f * h) * g = f * (h * g)$
- **Identity element:** $\delta * h = h * \delta = h$
- **Linear (superposition principle):**

$$f * (h_1 + h_2) = f * h_1 + f * h_2$$

$$(f + g) * h = f * h + g * h$$

$$(\alpha f) * h = \alpha(f * h)$$

- Mathematicians say: $(\{f : \mathbb{Z}^2 \rightarrow \mathbb{R}\}, +, *)$ is an “unital commutative algebra over \mathbb{R} ” (you may want to remember that to impress your friends).
- **Shift-invariant:**

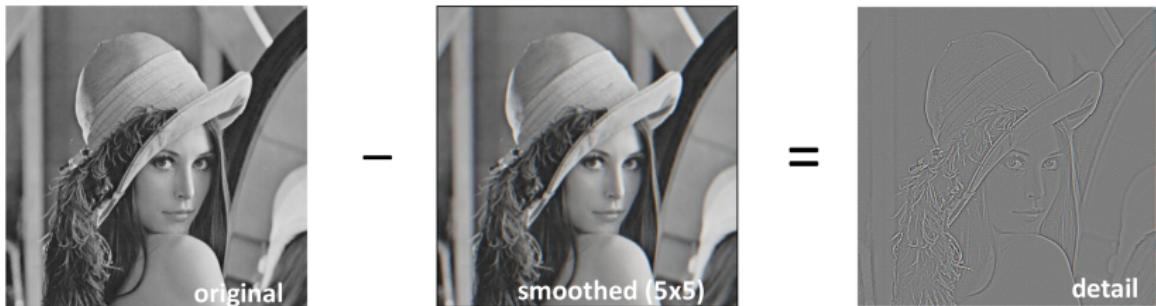
$$(\text{shifted } f) * h = (f * h) \text{ shifted.}$$

This is correct because shifting an image can be expressed by a convolution with a shifted impulse.

We have seen smoothing ... what about sharpening?

Sharpening idea

- What does blurring take away?



- Let's add it back!

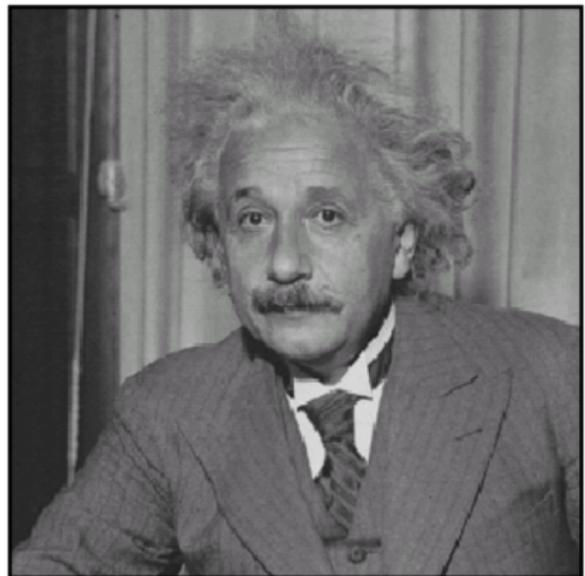


Sharpening filter

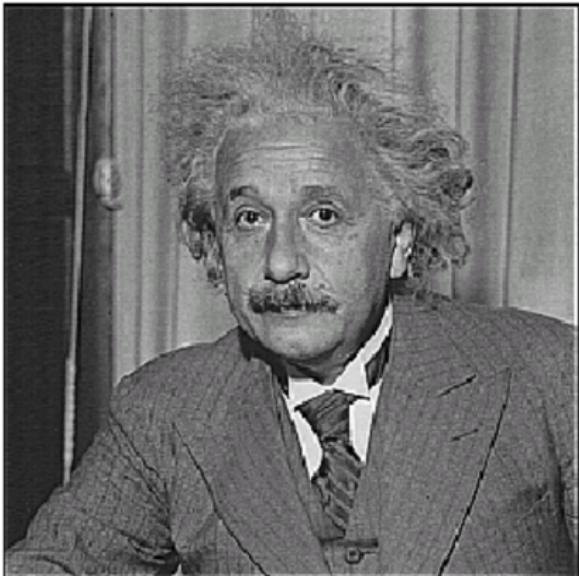
$$\text{Original} \quad \ast \left(\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix} - \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} \right) = \text{Result}$$

Sharpening filter
(accentuates edges)

Sharpening filter: results

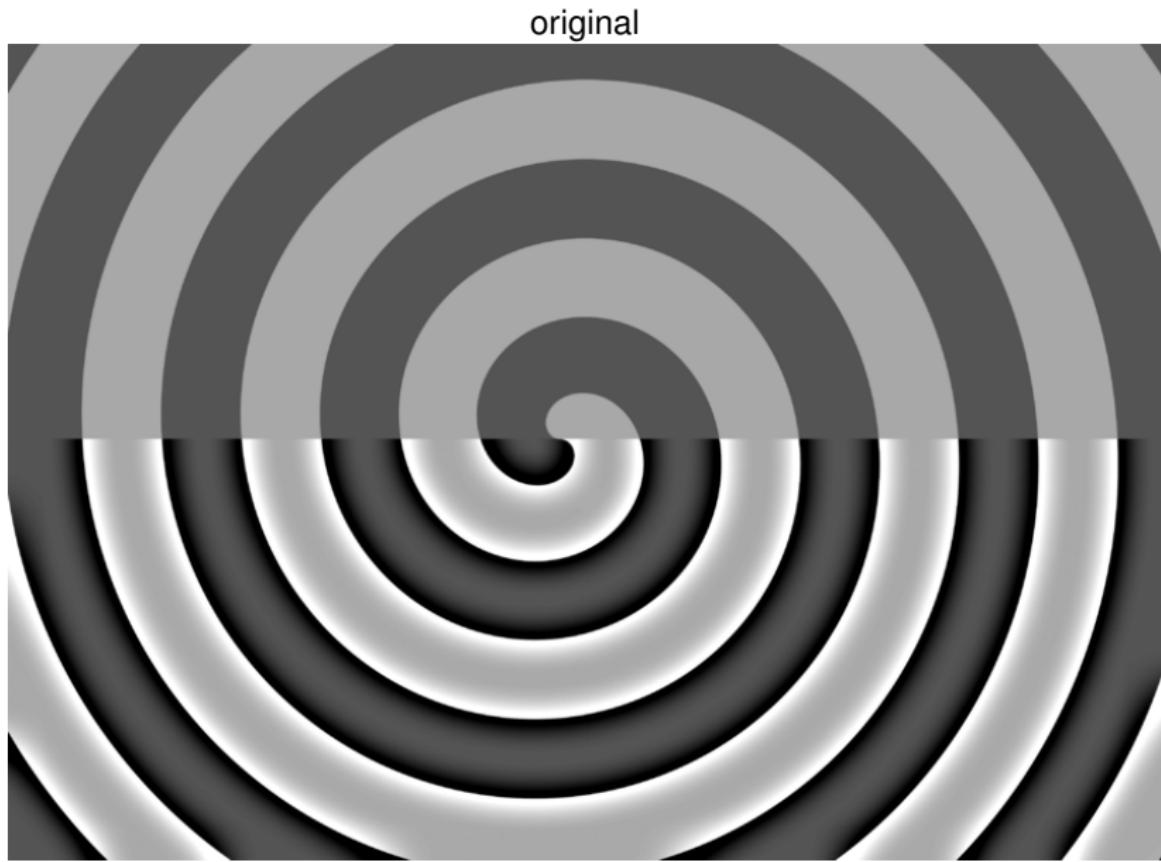


input



sharpened

Sharpening filter: results



Sharpening noisy images ... bad idea



Input with Gaussian noise $\sigma = 0.25f_{max}$



Result of sharpening, $\alpha = 3.0$

Note: noise is “high-frequency” content, i.e. detail.

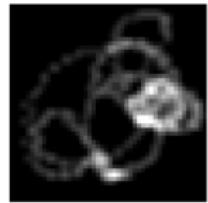
Convolution frequently occurs in optical systems

“Optical” convolution

- Camera shake or motion blur



=



“Optical” convolution

- Camera shake or motion blur



=



- Blur in out-of-focus regions of an image
(related to shape of the aperture, see later chapter).



Overview

1 Image types and discretization

2 Image noise

3 Correlation and Convolution

Noise reduction and moving average

The correlation

Gaussian filter

The convolution

Other filter kernels

4 Non-linear filters and denoising

Bilateral filter

Non-local means

Median filter

5 Summary

that was all nice and well, but remember the original goal -
what is the actual denoising performance of the filters?

Denoising Gaussian noise with a Gaussian filter



Gaussian additive noise, $\sigma = 5\% f_{max}$



Gaussian filter, $\sigma = 1.0$

Does not look too good subjectively - no preservation of edges.
But what is the best choice of parameter?

How to measure denoising quality? Two common measures.

Note: we often use the notation $(f_n)_{n=1,\dots,N}$ for all pixels in an image, if it is not important to distinguish the dimensions.

Mean square error (MSE)

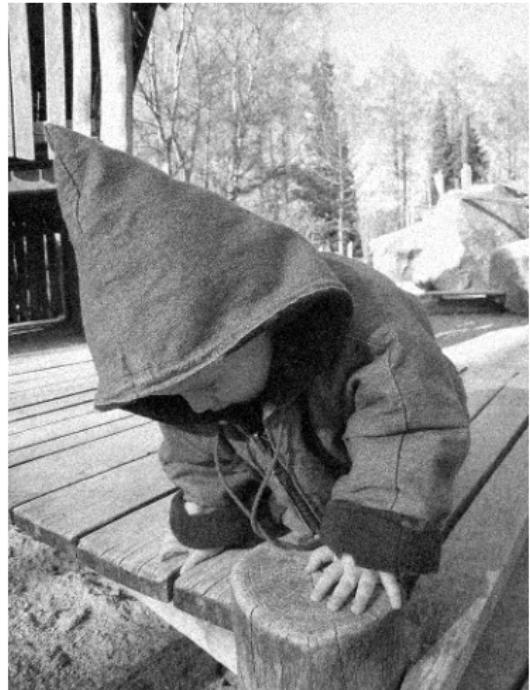
$$MSE(f, g) = \frac{1}{N} \sum_{n=1}^N (f_n - g_n)^2.$$

Peak signal-to-noise ratio (PSNR). For a maximum greyscale value of V ,

$$PSNR(f, g) = 10 \log_{10} \left(\frac{V^2}{MSE(f, g)} \right) \quad \leftarrow \frac{\text{peak signal}^2}{\text{noise}^2}$$

Note: both do not correspond very well to human perception - check out "structural similarity" (SSIM) for something better.

Denoising with Gaussian filter - best result



Gaussian additive noise, $\sigma = 5\%f_{max}$, PSNR=27.28

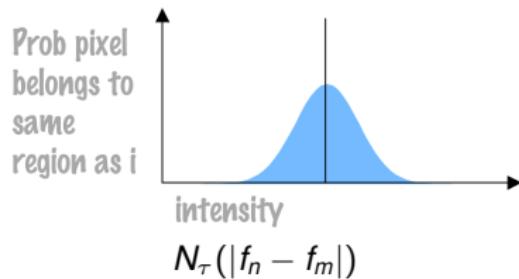
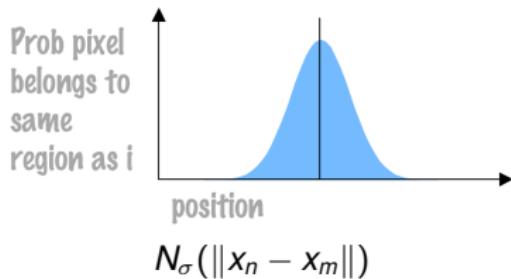


Gaussian filter, $\sigma = 0.5$, PSNR=29.18

The problem are the edges - how can we do better?

The bilateral filter

- **Key idea:** To preserve edges, average only over pixels which are nearby **and** have similar color.
- Two weights, one according to distance and one according to intensity, both Gaussian:



- **Note on notation:** x_n, x_m are the locations of the pixels, $1 \leq n, m \leq N$, and f_n, f_m the corresponding intensities. σ, τ are the standard deviations for the two Gaussian distributions N_σ, N_τ .

The bilateral filter

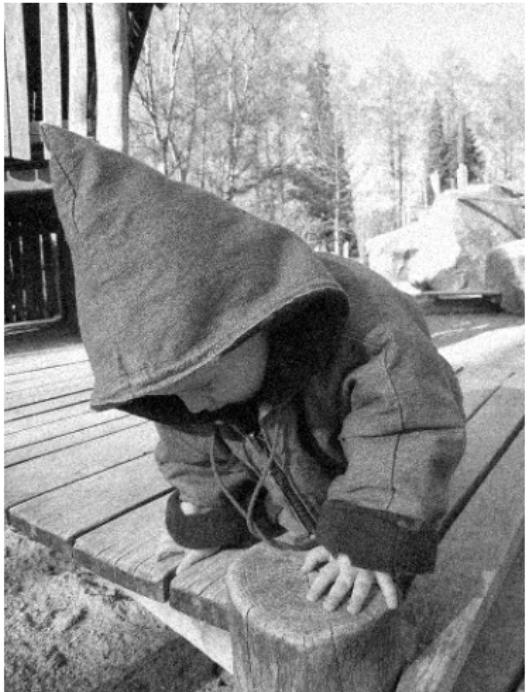
- For parameters $\sigma, \tau > 0$ and a greyscale image f , the bilateral filter in the n th pixel is defined as

$$g(n) = \frac{1}{K_n} \sum_{m=1}^N f_m N_\sigma(\|x_n - x_m\|) N_\tau(|f_n - f_m|)$$

with $K_n = \sum_{m=1}^N N_\sigma(\|x_n - x_m\|) N_\tau(|f_n - f_m|).$

- Note: Although the sum runs over all pixels, in practice, one will only compute it over a small window with a distance of about three times σ around x_n .
- Note: for $\tau \rightarrow \infty$, the result approximates a Gaussian filter with standard deviation σ .

Denoising with bilateral filter - best result in terms of PSNR



Gaussian additive noise, $\sigma = 5\%f_{max}$, PSNR=27.28



Bilateral filter, $\sigma = 1, 0.1$, PSNR=30.20

Compare to Gaussian filter - best result in terms of PSNR



Gaussian additive noise, $\sigma = 5\% f_{max}$, PSNR=27.28



Gaussian filter, $\sigma = 0.5$, PSNR=29.18

Non-local means

- **Observation:** images have similar patches all over the place
- **Idea:** When averaging, put most weight on the most similar patches
- **In formulas:**
 - around each pixel n , define a $k \times k$ window of pixels \mathcal{W}_n , which is written as a vector of pixel indices.
 - define a distance $d_{n,m}$ between pixels as

$$d_{n,m}^2 := \sum_{s=1}^{k^2} (f(\mathcal{W}_n(s)) - f(\mathcal{W}_m(s)))^2.$$

- kernel weights are then given by

$$k_{n,m} := \exp\left(-\frac{d_{n,m}^2}{2\sigma^2}\right),$$

where $\sigma > 0$ is a parameter.

[Baudes et al. 2005]

- As usual, with all kernel weights defined, the n th pixel in the filtered image g is then defined as

$$g(n) := \frac{1}{\sum_{m=1}^N k_{n,m}} \sum_{m=1}^N k_{n,m} f_m.$$

- Note: for practical reasons, the sums above are usually also restricted to a window around the n th pixel.
- Sometimes, the process is iterated to get a better estimate with already pre-denoised patches.
- However, in any case, this filter is usually very, very slow.

Denoising with NL means - best result in terms of PSNR



Original



NL means, $\sigma = 0.4$, PSNR=30.12

Compare to bilateral filter - best result in terms of PSNR



Original



Bilateral filter, $\sigma = 1, 0.1$, PSNR=30.20

Compare to Gaussian filter - best result in terms of PSNR



Original



Gaussian filter, $\sigma = 0.5$, PSNR=29.18

So far, we have only checked Gaussian noise -
what about other types of noise?

Denoising salt and pepper noise



Salt and pepper noise, $p = 0.2$



Local averaging filter, window size 5×5

Denoising salt and pepper noise



Salt and pepper noise, $p = 0.2$



Local averaging filter, window size 5×5

- Not zero mean locally
- Averaging biases the results locally

Median filter

- Very good for removing outliers
- For each pixel n , compute the vector \mathbf{f} consisting of k^2 grayscale values

$$\mathbf{f} = (f_{n,1}, \dots, f_{n,k^2})$$

in a local window of size $k \times k$ around n .

- The filtered result $g(n)$ at that location is defined as the **median value** of this vector. Informal definition: sort all elements in the vector by size and take the middle element of the result (i.e. the one at position $(k^2 + 1)/2$ if k is odd as usual).
- The underlying idea is a piecewise constant image model, we will skip the details here.

Median filter: results



Salt and pepper noise, $p = 0.2$

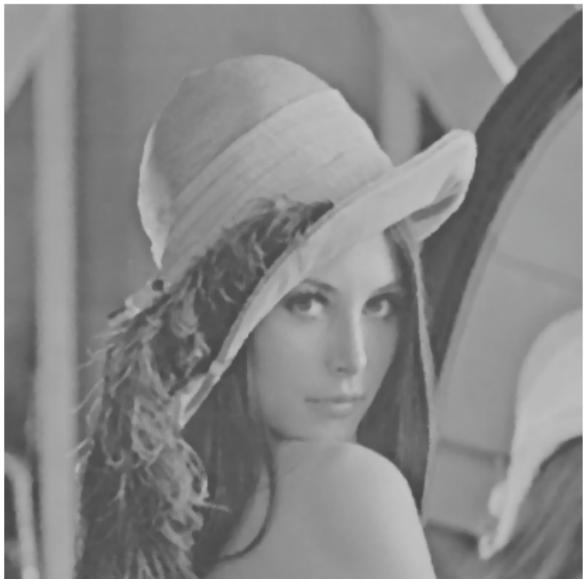


Median 3×3

Median filter: results



Salt and pepper noise, $p = 0.2$



Median 5×5

Median filter: results



Median 3×3 closeup



Median 5×5 closeup

Iterated median filter



Median 3×3 applied twice



Median 5×5 closeup

Overview

1 Image types and discretization

2 Image noise

3 Correlation and Convolution

Noise reduction and moving average

The correlation

Gaussian filter

The convolution

Other filter kernels

4 Non-linear filters and denoising

Bilateral filter

Non-local means

Median filter

5 Summary

Summary

- In image analysis, **convolution** is by far the most important filter.
- Convolution is linear and shift-invariant.
- For the application of image denoising, the best filters are non-linear and non-local.
- Know your noise, and you will be able to select the appropriate filter for denoising - ignore what noise you have, and failure to remove it properly will be a likely result.
- One can get better at denoising for example by learning how images should look like from training data - this is however seminar material.
- We will explore a few of the manifold applications of convolution in the next chapter, when we turn towards analyzing image derivatives and local features.